

CHAPTER III

Navigation

Navigation is the task of generating a series of commands that allow a robot to transit from its current pose to some goal pose. Within this task are a number of subtasks such as path planning, environment sensing, obstacle avoidance and others. As a robot transits through space from its current to goal pose it necessarily visits a series of intermediate poses that are a function of its kinematic constraints and the constraints of the environment it transits through. Collectively this set of poses as a time sequence is known as a path. The set of all possible poses is known as the task space and all kinematic feasible configurations is known as the configuration space. Typically the environment a robot traverses is not free space but also contains obstacles which constrain the paths which can be taken to the goal. While in some controlled environments the location of obstacles can be known *a priori* and supplied to the robot in the form of a map, often the robot will have the task of creating a map which describes the location of these obstacles. In some situations “global” sensors can view the entire environment which the robot will operate and can capture all of the constraints while in others (especially in the case of mobile robots) only “local” sensors mounted to the robot are available to integrate environmental limitations to the robot’s path as they are encountered. Even more challenging is the case where the environmental obstacles are not static but are moving adding a time varying element to the obstacle avoidance problem.

Algorithms that solve the navigation problem are sometimes broadly classified into four major categories (though there are many other categories which these algorithms could be grouped into). The first two, known as global or local, speak to the amount of spatial data or time horizon considered when planning a solution. If only a short time horizon or occlusions in the immediate spatial region are accounted for then the algorithm is considered local, otherwise it is thought of as global. The other two deal with how the spaces or paths are represented, continuously or discretely. If the algorithm

breaks up these spaces into finite resolution pieces then it is discrete. Continuous algorithms choose some sort of parameterization to the spatial or environmental constraints using a model. Often algorithms are some mixture of all of these four categories.

The algorithm used for navigation in this thesis is the GODZILA navigation algorithm, which is a form of local continuous variety based on the potential fields navigation strategy. While in many environments GODZILA solves the navigation problem, it can also be used in conjunction with global algorithms that can plan way points to the goal without having to consider the detailed local environment.

Something to consider is how different local sensors can affect the choice of navigation algorithm. The Nao humanoid platform used in this thesis is equipped with two sonar sensors for obstacle detection. Details on these can be read in Chapter II. These sensors have a broad angular range and cannot provide information about where within this cone obstacles are located. This allows the robot to avoid colliding with objects directly in front of it but restricts the available paths in the case where the width of the space between the obstacles approaches the width of the robot. Conversely, scanning laser rangefinders (such as the one mounted to the Nao and described in Chapter II) have a very high degree of angular resolution and provide hundreds of range measurements. This allows for a more accurate description of environmental occlusions allowing more paths to be considered for traversal.

The notion of what objects in the environment occlude a path and which do not can be viewed as a function of the gait (or mode of locomotion) used by the robot. If the robot is restricted to a plane, such as the case in many ground vehicles, then an object of similar size to the vehicle could prevent traversal through that position. If the robot can also fly then that object may not present as an impediment to precession towards the goal. While still not transiting through that exact position in 3D space, the 2D projection of the path onto the plane will appear to have moved through the obstacle due to the transformation of the obstacle set according to the different motion modality, making a 2D planner applicable to the problem. In Chapter ?? a crawling gait for the Nao robot is considered which allows it to go under occlusions that it would otherwise need to plan around.

3.1 Algorithm Classifications

Here some of the different notions common to navigation algorithms are explored. These classifications are sometimes useful for the navigation engineer to think about as

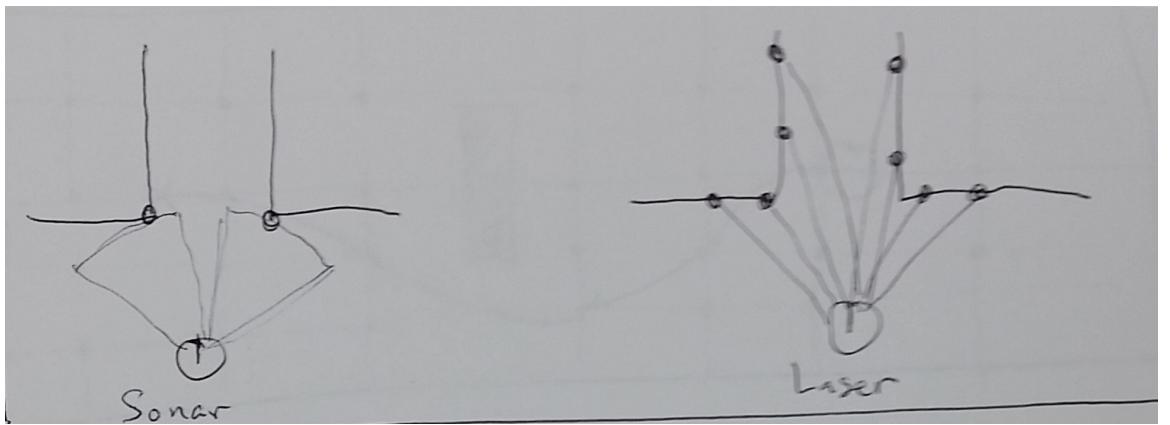


Figure 3: Sonar sensors are typically preferred when needing to detect the presence of an obstacle without needing precise position information. The location of a chair leg in the path of a robot would be detected by sonar while not knowing its position in the sonar cone. Laser rangefinders have very narrow **angular resolution**, typically requiring many of them to be effective at describing an environment. If many measurements are available at high angular **resolution** then they can be very effective at describing the width of narrow apertures.

tools available to solve the navigation problem. There are certainly other classifications and subclasses within these but this is intended as a sort of brief overview.

3.1.1 Global

One approach to the navigation problem is to consider the entire task and configuration space of a problem when generating a solution. For example, the task space of a robot arm might be the pose of the end effector and all of the occlusions within that space. The configuration space would be the set of all valid joint angles **the** arm can attain. Using **this** an algorithm can compute a time sequence of poses or movement commands that the robot should execute in order to bring it from its current pose to the final pose. While this approach is ideal in the sense of generating a solution to the original problem, it has several practical disadvantages. One requirement of the algorithm is having a complete description of the task space. In many cases global knowledge of occlusions while planning is not available meaning that when new obstructions are **observed** the algorithm needs to replan the path. The other commonly encountered problem with these algorithms **is the** magnitude of the space to be planned through is so large that computing a solution cannot be done in real-time.

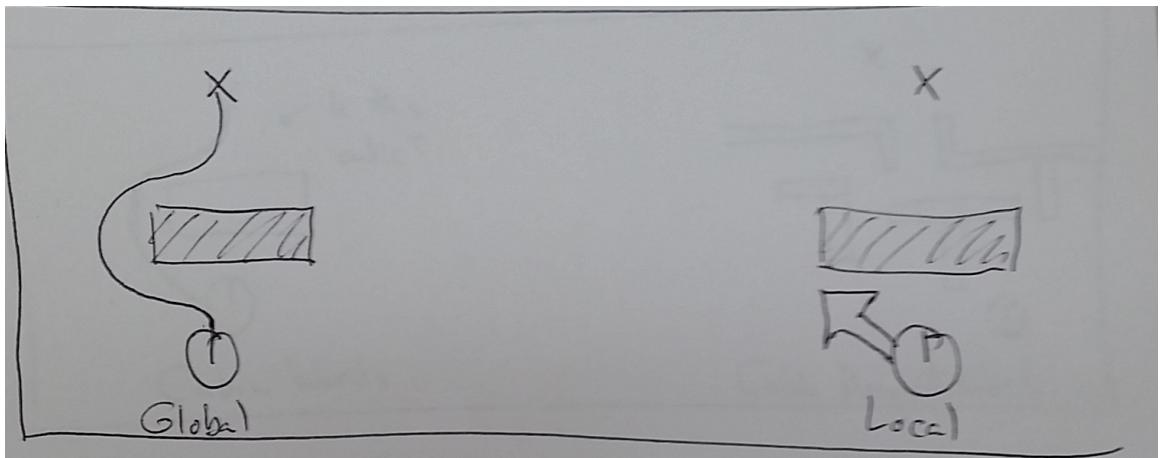


Figure 4: In global approaches an entire path is planned to the goal before commanding the robot. In local schemes the robot is given commands calculated as a function of its position and the environment and no one path is typically precomputed.

3.1.2 Local

A different approach to the navigation problem is generating commands based only on the current information available or a short history about the environment. With the robot arm example, there might be a sensor such as a camera that can detect objects that are obstructing the end effector from reaching the goal and instructs the arm to move around it. Such algorithms usually have the advantage of being quick to compute as compared to their global counterparts because they only have to consider a small subset of the task and configuration space. Their main disadvantage is that since they only use such a limited amount of information they can become trapped in local optima that do not allow the robot to achieve the desired pose.

3.1.3 Discrete

Orthogonal to the ideas of local and global path planning which deals with the scope of the time and space being considered when generating navigation commands is how these spaces are represented. One way to represent the space is to break it up into discrete representations and then plan through that. The task space for example could be divided into uniform spatial regions which an algorithm can consider visiting when planning a path. Alternatively, as with the robot arm example, a finite set of movement commands can be considered and iterated through while generating a solution. Graph-based search algorithms such as Dijkstra's algorithm or A* are examples

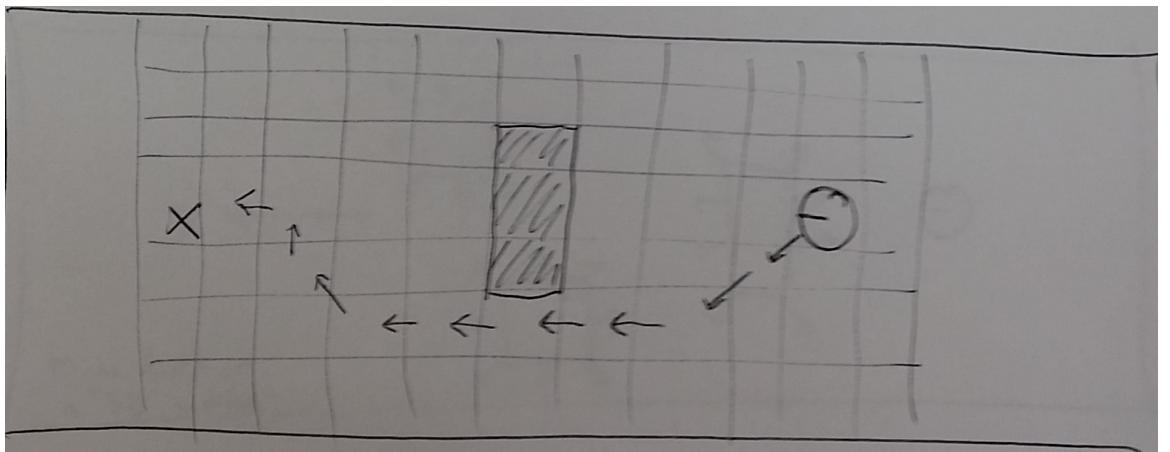


Figure 5: The environment here is being represented as occlusions in a discrete grid of locations. This allows the robot to plan through successive locations to the goal using graph-based techniques.

of discrete solvers. An advantage to this is paths with complex shapes can be generated to accommodate difficult constraints. A problem with this approach is how finely to resolve the task or configuration spaces. Coarser discretization can allow a solution to be computed rapidly but miss more optimal solutions.

3.1.4 Continuous

Continuous spatial representation avoids the problem of choosing a discretization resolution. The movement commands or paths are planned in a continuum allowing paths to take on intermediate values not available at a given discrete resolution. Continuous representation instead has a different problem of paths needing to take on particular solution forms. A path through the task space might be represented as a cubic spline from the current pose to the goal pose. This path has access to the entire task space but may not be able to construct a path under complicated environments with many obstacles. Often designers instead of using a single cubic will use a series of cubics, the termination of one being the starting position of the next, until the goal is reached. One example of such a planner is known as a “maneuver-based” planner.

3.1.5 Composite Approaches

In order to combine the strengths and combat the weakness the above approaches are often combined to form a more robust navigation solution. Global plans can be generated as a series of intermediate goals or way points for local planners to navigate

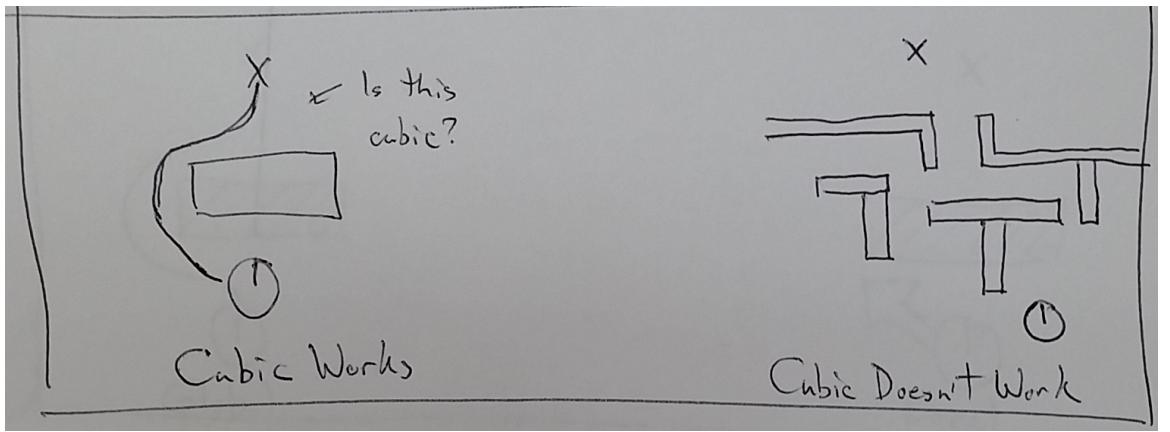


Figure 6: This figure shows two sample environments. On the **left** the robot can follow an **trajectory** that takes a cubic form and reaches the goal location. On the **right** a more complicated environment **where** a single cubic can not describe an appropriate path.

towards. These way points can be planned on a coarse discrete grid that is quick to compute while low dimensional continuous trajectories are plotted between them. Lattice planners are an example of such a composite algorithm.

3.2 Potential Fields

Potential fields algorithms are continuous local planners. They work on the concept of modeling the robot as a sort of charged particle (like an electron) and obstacles in the environment produce a repulsive field pushing the robot away from them. The goal location in turn acts as an attractive force, pulling the robot towards it. As the robot moves through the environment it detects objects and generates motion commands based on their range and direction.

Given a vector x_r which describes the position of the robot in the task space and the position of obstacles m_i in the set of all obstacles M , a force vector f_i can be generated for each obstacle that describes the magnitude and direction of the repulsive force applied to the robot by the obstacle. The force direction is in the opposite direction of the relative bearing of the obstacle to the robot with magnitude being inversely proportional to the relative range of the obstacle to the robot. Equation (3.2) shows an example of force vector generation.

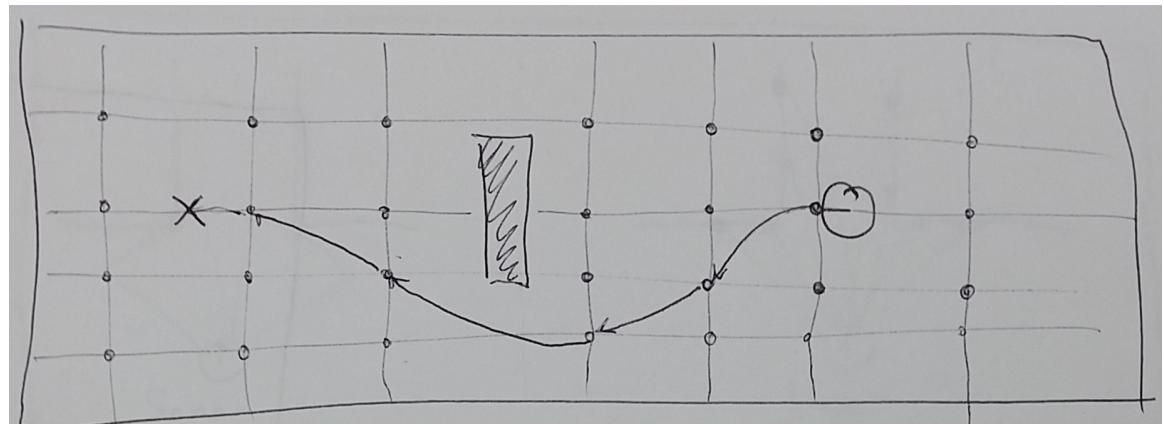


Figure 7: A lattice planner is an example of a composite discrete and continuous path planner. The planner first plans through a discrete grid and then designs trajectories that bring the robot from one point in the grid to the next.

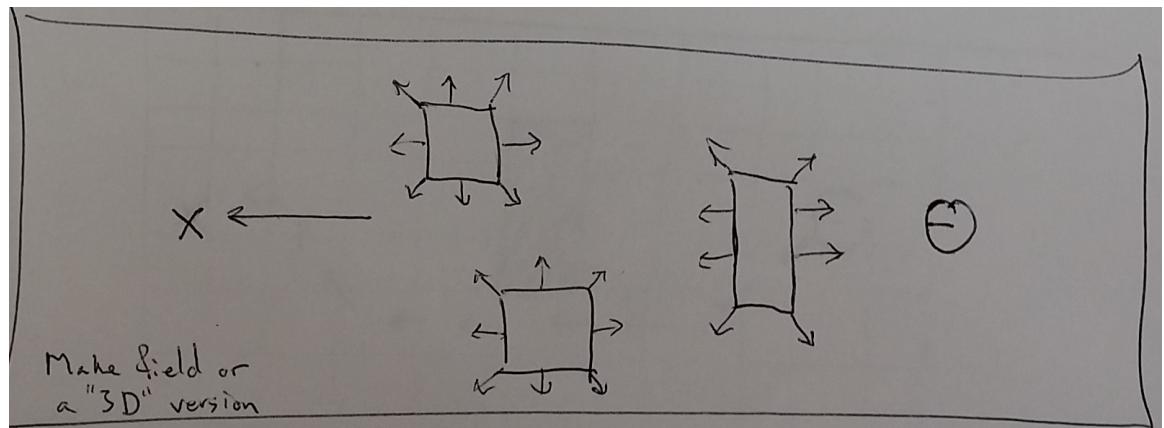


Figure 8: Visualization of potential fields idea. Objects in the environment emit a field that repels the robot with the goal emitting an attractive field. This can be visualized as a 2D vector field.

$$\|f_i\| = \frac{-c_i}{\|m_i - x_r\|^\alpha} \quad (3.1)$$

$$\angle f_i = \angle(m_i - x_r) \quad (3.2)$$

$m_i - x_r$ is the relative location of the obstacle to the robot and c_i is a positive scalar used to tune force magnitude. In some implementations c_i can be a function of the relative bearing, for example, decreasing the repulsion affect if the obstacle is not in the direction of the goal. The parameter α in general can be any positive scalar but is often picked to be 2.

For the goal seeking behavior, the equations are similar, with the direction of the force being in the direction of the goal, rather than away from it.

$$\|f_g\| = \frac{c_g}{\|x_g - x_r\|^\gamma} \quad (3.3)$$

$$\angle f_g = \angle(x_g - x_r) \quad (3.4)$$

$x_g - x_r$ in equation (3.4) is the relative location of the goal to the robot and c_g is some positive tuning scalar. γ is some positive scalar, typically 2.

The sum of these forces is used to produce a force vector f_r used to command the robot away from obstacles and towards the goal location.

$$f_r = f_g + \sum_{i=1}^{\|M\|} f_i \quad (3.5)$$

3.3 GODZILA

GODZILA (Game-theoretic Optimal Deformable Zone with Inertia and Local Approach) [?] is a local continuous navigation algorithm based on the potential fields idea. In contrast to some potential fields formulations which work on the range and bearing of objects in the environment, GODZILA uses the sensor information about occlusion locations directly in order to formulate navigation commands. It is a memoryless algorithm and does not attempt to build a map of the environment making it very lightweight in terms of both computation and memory. It can be implemented on a variety of vehicles and is suited for the cases where a low computation microcontroller is utilized.

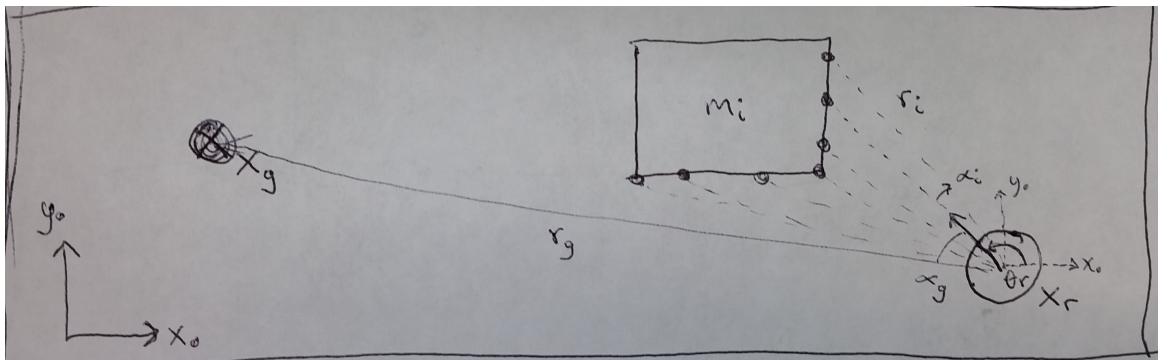


Figure 9: Illustration of all the GODZILA variables. Range and bearing of the goal to the robot are requisite as well as being able to take range measurements to occlusions and knowing their relative bearings.

It has three components which are an optimization cost, a straight line planner, and a stochastic local minima escape strategy.

Figure 9 diagrams the key variables in the formulation.

Taking a 2D example of the application of the GODZILA algorithm, we have the following terms:

x_0 is the x-axis of the inertial frame.

y_0 is the y-axis of the inertial frame.

x_r is the (x,y) location of the robot in the inertial frame.

α'_r is the new heading the robot should be commanded to in the robot frame.

x_g is the (x,y) location of the goal in the inertial frame.

r_g is the range to the robot relative to the position of the robot.

α_g is the bearing to the goal location relative to the orientation of the robot.

z_i is the i^{th} sensor measurement made by the robot which includes range and bearing to an occlusion.

r_i is the i^{th} range measurement to an occlusion relative to the position of the robot.

α_i is the bearing angle to the i^{th} sensor measurement.

In this formulation it is assumed that there is some mechanism to allow the robot to get an estimate of the relative range and bearing to the goal location r_g and α_g and that the robot has some sensor that can provide range and bearing to occlusions in the environment.

3.3.1 Optimization Formulation

The formulation of the GODZILA algorithm is in terms of an optimization cost. The goal of the algorithm is to bring the robot towards the goal while preventing it from colliding with obstacles. With this intuition the cost function has three components:

1. $J_1(\alpha'_r)$ which rewards goal seeking by penalizing directions other than those towards the goal.
2. $J_2(\alpha'_r)$ which penalizes directions towards occlusions.
3. $J_3(\alpha'_r)$ which dampens oscillations by penalizing changes in heading.

The sum of these terms produces the cost function to be minimized.

$$\min_{\alpha'_r} \sum_{i=1}^3 J_i(\alpha'_r) \quad (3.6)$$

Where α'_r is the heading direction the robot should travel in to minimize the cost function. $J_1(\alpha'_r)$ takes on the form of:

$$J_1(\alpha'_r) = g_{11}(r_g)f_{11}(\|\alpha'_r - \alpha_g\|) \quad (3.7)$$

g_{11} is a class- \mathcal{L} function meaning that it is continuous and monotonically non-increasing which maps $[0, \infty) \mapsto [0, \infty)$ and f_{11} is a class- \mathcal{K} function meaning it is monotonically non-decreasing which maps $[0, \infty) \mapsto [0, \infty)$. The intuition here is that that f_{11} produces a higher cost in proportion to how much the direction command α'_r differs from the goal direction. g_{11} conversely reduces the cost if the range to the goal is large. The idea being that the farther from the goal the robot is, then being oriented towards the goal becomes less of a priority.

$J_2(\alpha'_r)$ takes on the form of:

$$J_2(\alpha'_r) = J_{2_{\mathcal{I}_1}}(\alpha'_r) - J_{2_{\mathcal{I}_2}}(\alpha'_r) \quad (3.8)$$

$$J_{2_{\mathcal{I}_1}}(\alpha'_r) = \sum_{i \in \mathcal{I}_1} g_{21}(r_i) \left[g_{22}(\|\alpha_g - \alpha_i\|) + g_{23}(\|\alpha_i\|) \right] g_{24}(\|\alpha'_r - \alpha_i\|) \quad (3.9)$$

$$J_{2_{\mathcal{I}_2}}(\alpha'_r) = \sum_{i \in \mathcal{I}_2} f_{21}(r_i) g_{25}(\|\alpha_g - \alpha_i\|) g_{26}(\|\alpha'_r - \alpha_i\|) \quad (3.10)$$

Equation (3.8) has two components to it, separated by the membership of the sensor measurements z_i to one of two sets. The first is the set of all range measurements r_i whose value falls below some positive scalar r_c are put into a set \mathcal{I}_1 . The second set is the remaining range measurements, which are necessarily greater than r_c , placed into set \mathcal{I}_2 . This threshold distance r_c is picked such that if the occlusion is farther away than it, the robot does not need to be concerned with avoiding it. Equation (3.9) is concerned with the occlusions that need to be avoided, and has four class- \mathcal{L} functions $g_{21}, g_{22}, g_{23}, g_{24}$ associated with it. The first three g_{21}, g_{22}, g_{23} act as gains that the control variable α'_r in function g_{24} must attenuate. g_{21} says that if the range r_i to the occlusion is small, then the cost of orienting in that direction is high. g_{22} amplifies this effect if the direction to the occlusion α_i is in the same direction of the goal α_g . g_{23} also amplifies g_{21} if the direction of the occlusion is similar to the heading direction of the robot. Equation (3.10) deals with the case where the occlusions can be approached. It has one class- \mathcal{K} function f_{21} and two class- \mathcal{L} functions g_{25} and g_{26} . f_{21} and g_{25} act as gains on the controlled function g_{26} . f_{21} promotes moving in the direction of ranges that are large and g_{25} promotes moving in the direction of occlusions that are in a similar direction to the goal direction.

The final term $J_3(\alpha'_r)$ is a class- \mathcal{K} function that penalizes changes in direction.

$$J_3(\alpha'_r) = f_{31}(\|\alpha'_r\|) \quad (3.11)$$

This term is present to prevent high frequency oscillations and is analogous to giving the robot a physical inertia. All of the functions $f_{11}, f_{21}, f_{31}, g_{11}, g_{21}, g_{22}, g_{23}, g_{24}, g_{25}, g_{26}$ in the optimization are tunable by the designer but it was noted in [?] that in the case where the controlled functions $f_{11}, g_{24}, g_{26}, f_{31}$ are chosen to be quadratic. In this case the optimization problem can be solved in closed form producing the solution seen in equation (3.12).

$$\alpha'_r = \frac{\alpha}{\|\alpha\|} \quad (3.12)$$

$$\alpha = \alpha_1 + \alpha_2 + \alpha_3 + \alpha_4 \quad (3.13)$$

The four terms in equation (3.13) correspond to the solutions to the four optimization terms $J_1(\alpha'_r), J_{2\mathcal{I}_1}(\alpha'_r), J_{2\mathcal{I}_2}(\alpha'_r), J_3(\alpha'_r)$. They are expanded in equations (3.14)-(3.17).

$$\alpha_1 = \bar{f}_{11}g_{11}(r_g)\alpha_g \quad (3.14)$$

$$\alpha_2 = -\bar{g}_{24} \sum_{i \in \mathcal{I}_1} g_{21}(r_i) \left[g_{22}(\|\alpha_g - \alpha_i\|) + g_{23}(\|\alpha_i\|) \right] \alpha_i \quad (3.15)$$

$$\alpha_3 = \bar{g}_{26} \sum_{i \in \mathcal{I}_2} f_{21}(r_i) g_{25}(\|\alpha_g - \alpha_i\|) \alpha_i \quad (3.16)$$

$$\alpha_4 = \bar{f}_{31}\alpha_r \quad (3.17)$$

The solutions to the equations reuse functions $f_{21}, g_{11}, g_{21}, g_{22}, g_{23}, g_{25}$ as they are not functions of α'_r and produce constant versions of $f_{11}, f_{31}, g_{24}, g_{26}$ as $\bar{f}_{11}, \bar{f}_{31}, \bar{g}_{24}, \bar{g}_{26}$ in the direction of their respective directions $\alpha_g, \alpha_i \forall i \in \mathcal{I}_1, \alpha_i \forall i \in \mathcal{I}_2, \alpha_r$ where $\alpha_r = [1, 0]^T$ representing the current heading of the robot in the vehicle frame.

The resultant α'_r is typically commanded as an angular rate $\dot{\alpha}'_r$ according to the angular velocity bandwidth of the vehicle.

Finally, the linear velocity is still to be commanded. A good pick for this function is something of the form in equation (3.18).

$$v_r = f_v(\min(R))g_v(\dot{\alpha}'_r) \quad (3.18)$$

Where R is the set of all range measurements, f_v is a class- \mathcal{K} function and g_v is a class- \mathcal{L} function. This reduces the speed of the vehicle when it is in close proximity to occlusions or if it is commanded to a high angular rate. Slower speeds near obstacles makes it less likely that the vehicle will collide with them. Slower speeds when commanded to high angular rates helps reduce the distance the robot travels in a previously commanded direction before completing a new turn command.

3.3.2 Straight Line Planner

If at some point during the navigation of the vehicle to the goal the straight line path from the robot to goal becomes unobstructed, it follows that the robot should

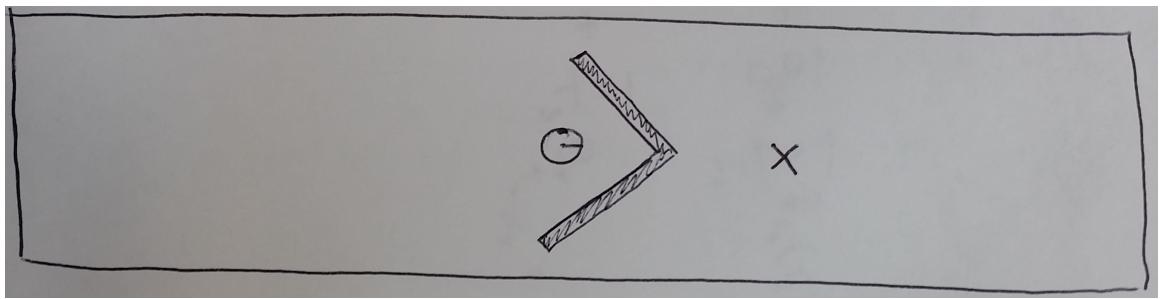


Figure 10: It's a trap!

proceed directly to the goal according to this path. This is sensible, but can bring the robot in close proximity to obstacles. It is therefore desirable to preserve the obstacle avoidance properties of the optimization procedure from 3.3.1. This can be achieved by prescribing intermediate goal locations according to (3.19). This procedure is especially useful in situations where the width of the aperture between occlusions to the goal approaches the width of the robot where the obstacle avoidance behavior might deter traversal to the goal.

$$\hat{x}_g(t) = (1 - \lambda(t - t_0))x_r + \lambda(t - t_0)x_g \quad (3.19)$$

Where $\lambda(\tau)$ is a monotonically increasing function that maps $[0, T_f] \mapsto [0, 1]$ with T_f being some amount of time allowed for the vehicle to reach the goal location. The length of time the planner is active is T_f . If the robot has not reached the goal in this time but again has a straight line path to the goal the planner is allowed to restart.

3.3.3 Escape Strategy

While the inertial term provided by (3.11) is intended to reduce high frequency oscillations and backtracking that might trap the robot in small corners it is not enough to account for degenerate cases such as those shown in Figure 10.

In such scenarios an effective strategy to escape such traps is to randomly assign an alternate goal position x_{rand} that will allow the robot to escape the local minima and proceed to the goal. The distance to x_{rand} from the robot is typically selected to be smaller than the distance to the nearest occlusion $\min(R)$. After some fixed time the goal position is set back to x_g . If the trap condition is detected multiple times then the next escape procedure is repeated that many times before being restored to x_g . One possible method to detect a trap condition is to integrate a sliding window of vehicle linear velocities v_r and if those values integrate to a position close to zero then the robot

is trapped. It is shown in [?] that with all of these mechanisms the position of the robot x_r will converge to x_g in finite time with probability 1.